

Chapter 4 - HiveQL

Exercise

1. How many rows are in the Movies table? How many rows are in the UserRatings table?
2. Using the Movies table, find the name of all movies released in 1990.
3. List the movieid of the 10 most rated films in the UserRatings table (i.e., which ones have the most ratings).
4. Use a join to list the names of the movies you found in step 3.
5. Do any movies have no ratings? Hint: use an outer join and IS NULL
6. What is the highest rated sci_fi movie?
7. What is the highest rated sci_fi movie that has at least 10 user ratings? Hint: use a subquery