



Pig Latin basics



In this chapter, you will learn

- What is Pig Latin?
- How to use the basic Pig Latin commands for filtering, grouping and sorting data
- Under the rules of case-sensitivity



What is Pig Latin?

- A data flow language composed of a series of statements
- Each statement operates on a relation, does some transformation and returns a relation



Pig Latin is a scripting language used to express data flows as a series of statements (also called operations or transformations). Pig will convert the statements into a series of MapReduce jobs. A Pig program can accomplish rather sophisticated data processing with only a few lines of code as opposed to writing 100s of lines of code in Java MapReduce.

Each Pig Latin statement takes a relation as input (except the LOAD statement, which creates the initial relation). The statement does some operation and returns a new relation as output. It is common to give each resulting relation a new alias.

Most statements in Pig Latin end in a semicolon (with the exception of some statements like the Hadoop filesystem commands).

An example - find the highly rated movies

```
ratings = LOAD 'u.data' AS (userid:int,
    itemid:int, rating:int, timestamp:int);

grp = GROUP ratings BY itemid;

avg_rating = FOREACH grp GENERATE group,
    AVG(ratings.rating) AS score;

good_movies = FILTER avg_rating BY score>4.5;
```



The u.data file contains 100000 ratings by 943 users on 1682 movies. Each user has rated at least 20 movies. The file contains a tab-separated list of : user id, item id, rating and timestamp.

The following script uses the u.data file to find the movies that received an average rating over 4.5:

-- Load the u.data file which contains 4 integers, tab-separated. Default PigStorage loader is used. "ratings" is the alias, or name, given to the result

```
ratings = LOAD 'u.data' AS (userid:int, itemid:int,
    rating:int, timestamp:int);
```

-- Group the movies in the relation "ratings" by their itemid

```
grp = GROUP ratings BY itemid;
```

-- Calculate an average score per group (i.e., per movie). Note: lower case "group" is a keyword that we'll talk about later.

```
avg_rating = FOREACH grp GENERATE group,
    AVG(ratings.rating) AS score;
```

-- Remove all movies that did not score better than 4.5

```
good_movies = FILTER avg_rating BY score > 4.5;
```

FILTER..BY

- Removes data from a relation if it does not match some criteria

- Syntax:

```
alias = FILTER alias BY expression;
```

- Examples:

```
b = FILTER a BY id == 42;
```

```
b = FILTER a BY name == 'Doug';
```

```
b = FILTER a BY $0 > 5 AND $1 == 'hi';
```



The `FILTER` operator will remove (or filter) records from a relation if they do not match some criteria. This could be a basic expression such as "does the id field equal 42?" or a compound expression that looks at multiple fields. The general syntax is:
`alias2 = FILTER alias1 BY expression;`

`alias1` is the incoming relation and `alias2` is the resulting (probably smaller) relation.

Here are several more examples:

-- Find all records in `a` that have value 42 for their id field:

```
b = FILTER a BY id == 42;
```

-- Find all records in `a` that have the exact string 'Doug' in their name:

```
b = FILTER a BY name == 'Doug';
```

-- Find all records in `a` that have the value 5 for their first field and string 'hi' for their second field:

```
b = FILTER a BY $0 > 5 AND $1 == 'hi';
```

FOREACH . . GENERATE

- FOREACH can do an operation on each record in a relation
- Syntax:

```
alias = FOREACH alias  
    GENERATE expression [,expression..];
```

- Example:

```
b = FOREACH a GENERATE $0;  
b = FOREACH a GENERATE f1, f3 + 42;  
b = FOREACH a GENERATE SIZE(f2);
```



FOREACH can apply an operation on each record of a relation. In other words, it iterates through the records. Some examples:

-- For each record in a, output just the first field (\$0):

```
b = FOREACH a GENERATE $0;
```

-- For each record in a, output the field named f1 and the field f3 plus 42.

```
b = FOREACH a GENERATE f1, f3 + 42;
```

Some uses of FOREACH:

- Remove or reorder some columns
- Apply operators (+, -, *, /, etc) to the data
- Apply functions to the data (covered later)
- Iterate through a grouped relation (covered later)

Eliminating duplicates

- Use `DISTINCT` to filter duplicate records

- Syntax:

```
alias = DISTINCT alias;
```

- Examples:

```
b = DISTINCT a;
```

```
names = FOREACH persons GENERATE name;
```

```
uniques = DISTINCT names;
```



The `DISTINCT` command eliminates duplicate records in a relation. This applies to the whole record - in order to be considered a duplicate, all fields must be equal.

For example, let's take this data:

Doug	Cutting
Doug	Jones

Loading the data and applying `DISTINCT` will not remove any rows because the records have different last names. However, `FOREACH...GENERATE` can be used to create a new relation with only certain fields and then `DISTINCT` can be applied:

```
persons = LOAD 'data' AS (fname:chararray,  
lname:chararray);  
names = FOREACH persons GENERATE fname;  
uniques = DISTINCT fname;
```

GROUP . . BY

- Groups records with a similar value

- Syntax:

```
alias = GROUP alias BY expression;
```

- Examples:

```
b = GROUP a BY $0;
```

```
grpdb = GROUP movies BY release_date;
```



The `GROUP` keyword is used to group records in a relation into groups. Items in a group have a value in common (the group key). The basic syntax is:

```
alias2 = GROUP alias1 BY expression;
```

`alias1` is the incoming relation. `expression` is the criterion used to identify the groups. `alias2` is the resulting relation which contains one tuple per group. Each tuple looks like:

(group-key, bag-of-tuples-in-group)

Example

```
> cat data;
Doug      cat
Tom       dog
Mike      cat
Sarah     fish

> a = LOAD 'data' AS (name:chararray,
  pet:chararray);

> b = GROUP a BY pet;

> dump b;
(cat, {(Doug,cat), (Mike,cat)})
(dog, {(Tom,dog)})
(fish, {(Sarah,fish)})
```



Let's say the file had peoples' names and their favorite type of pet:

```
Doug      cat
Tom       dog
Mike      cat
Sarah     fish
```

We could load the data like this:

```
a = LOAD 'data' AS (name:chararray, pet:chararray);
```

Then group by the type of pet:

```
b = GROUP a BY pet;
```

The result would be 3 groups (cat, dog and fish):

```
(cat, {(Doug,cat), (Mike,cat)})
(dog, {(Tom,dog)})
(fish, {(Sarah,fish)})
```

Using the grouped results

- FOREACH works for grouped data too
- The grouped relation has a special field named "group"

```
a = LOAD 'data' AS
  (name:chararray,pet:chararray);
b = GROUP a BY pet;
c = FOREACH b GENERATE group, COUNT(a);
dump c;
(cat,2L)
(dog,1L)
(fish,1L)
```

implicit field name
given to the group key

To refer to a specific
field, use a.field

cloudera

Earlier we saw that FOREACH iterates through the records in a relation. This is useful for grouped results. Example:

```
a = LOAD 'data' AS (name:chararray, pet:chararray);
b = GROUP a BY pet;
-- For each record in b, output the group (pet) and the
number of records (a). An alias can be given to the
result of COUNT using AS
c = FOREACH b GENERATE group, COUNT(a) AS num;
DUMP c;
(cat,2L)
(dog,1L)
(fish,1L)
```

Relation `b` has a field named "group" which refers to the group key (the pet).

It is common to use aggregate functions, such as `COUNT()` on groups (covered later).

GROUP ..ALL

- Use GROUP ..ALL to put all records in a single group

- Syntax:

```
alias = GROUP alias ALL;
```

- Example:

```
a = LOAD 'data' AS  
  (name:chararray,pet:chararray);  
b = GROUP a ALL;  
c = FOREACH b GENERATE COUNT(a);
```



Sometimes all the records should be collected into a single group. This is usually done to calculate aggregate functions on ALL records. For example, we can count the number of records in the file using:

```
a = LOAD 'data' AS (name:chararray, pet:chararray);  
b = GROUP a ALL;  
c = FOREACH b GENERATE COUNT(a);
```

ORDER..BY

- Use ORDER..BY to sort the records in a relation
- Syntax:

```
alias = ORDER alias BY field [DESC];
```
- Examples:

```
b = ORDER a BY lastname;  
b = ORDER a BY lastname, firstname;  
b = ORDER a BY age DESC;
```



To sort records in a relation by a field or fields, use ORDER..BY. By default the sort order is **ascending**. Adding the keyword **DESC** will sort them descending.

Multiple fields can be specified, such as ORDER..BY lastname, firstname. This will sort by lastname, with a secondary sort by firstname (like the white pages in a phone book). Each field may choose ascending or descending order.

Pig can take advantage of multiple reducers (see PARALLEL keyword later). A special Partitioner is used to accomplish this.

Note: In the current version of Pig, the order by field must be a column and cannot contain expressions. This will change in a later version (Pig 0.9).

Watch out for non-typed data

```
> cat data
```

```
42
```

```
100
```

```
3
```

```
> a = LOAD 'data';
```

```
> b = ORDER a BY $0;
```

```
> DUMP b;
```

```
(100)
```

```
(3)
```

```
(42)
```



bytearrays order
by byte order, not
numerically

cloudera

Be careful when sorting data with an unspecified schema. If a schema was not given in the `LOAD`, the data is considered type `bytearray`. Sorting by byte order is not the same as sorting numerically. If you want to sort integers properly, make sure they are typed as integer.

Correct approach:

```
a = LOAD 'data' AS f1:int;
```

```
b = ORDER a BY f1;
```

```
DUMP b;
```

```
(3)
```

```
(42)
```

```
(100)
```

LIMIT

- Use `LIMIT` to reduce the number of output records

- Syntax:

```
alias = LIMIT alias n;
```

- Example:

```
b = LIMIT a 10;
```



The `LIMIT` keyword reduces the amount of records in a relation. This can be very useful and also makes Pig scripts more efficient by eliminating the amount of data that needs to be processed.

Unless an `ORDER BY` is also specified, the records returned from a `LIMIT` are somewhat random and may change from one execution to another.

Top-N queries

- Use `ORDER BY` and `LIMIT` together for top-n results

- Example:

```
ordered = ORDER items BY cost DESC;  
expensive = LIMIT ordered 10;
```



A "Top-N" query is looking for a certain number (N) of the top (i.e., greatest or least) items. For example, the *10 most expensive items* or the *100 worst movies of all time*.

This requires 2 steps:

1. Order the results by the criterion that makes sense
2. Use `LIMIT` to pull off the number of results you want to return

For example, the 10 least expensive items:

```
ordered = ORDER items BY cost;  
cheap = LIMIT ordered 10;
```

To ask for the 10 *most* expense, sort descending instead of ascending:

```
ordered = ORDER items BY cost DESC;  
expensive = LIMIT ordered 10;
```

Nested ordering

- ORDER BY can be applied within each group

- Example:

```
a = LOAD 'data' AS
    (name:chararray,pet:chararray);
b = GROUP a BY pet;
c = FOREACH b {
    ordered = ORDER a BY name DESC;
    GENERATE group, ordered;
}
```



Let's say the file had peoples' names and their favorite type of pet:

Doug	cat
Tom	dog
Mike	cat
Sarah	fish

```
a = LOAD 'data' AS (name:chararray, pet:chararray);
b = GROUP a BY pet;
-- For each group (cat, dog and fish), order the owners' names alphabetically
c = FOREACH b {
    ordered = ORDER a BY name;
    GENERATE group, ordered;
}
```

LIMIT can also be used in the FOREACH block.

Adding comments to a script

- Single-line comments:
`-- This is a comment`
- Multi-line comments:
`/*
 This is a longer
 comment.
*/`



Commenting code is a good way to add readability to your Pig scripts. There are two types of comments. Double hyphens are a single-line comment: from the hyphens, the rest of the line is ignored. Examples:

```
-- Load the data:  
a = LOAD..  
-- Find the records that match pattern X:  
b = FILTER..  
DUMP c; -- Print the results to the screen
```

C-style comments are also supported. They allow multiple lines to be ignored by the Pig parser:

```
/*  
    This is a longer  
    comment.  
*/
```

Case-sensitivity

Case-sensitive	Case-insensitive
Aliases (names of relations and fields), Functions (COUNT, AVG, PigStorage, etc.) String literals	Keywords (LOAD, USING, FILTER, ls, copyFromLocal, quit, etc.)



Pig Latin has mixed rules on case-sensitivity. Aliases and functions are case-sensitive, yet keywords and operators are not. Take this example:

```
ratings = LOAD 'u.data' AS (userid:int, itemid:int,
rating:int, timestamp:int);
grpds = GROUP ratings BY itemid;
avg_rating = FOREACH grpds GENERATE group,
AVG(ratings.rating) AS score;
good_movies = FILTER avg_rating BY score > 4.5;
```

In the above script, the aliases and the AVG function **are** case-sensitive. String comparisons such as `name = 'doug'` are also case-sensitive.

However, these are **not** case-sensitive:

```
LOAD, AS, int, GROUP, BY, FOREACH, GENERATE, FILTER
```

In this chapter, you have learned

- What is Pig Latin?
- How to use the basic Pig Latin commands for filtering, grouping and sorting data
- Under the rules of case-sensitivity

