



Advanced Pig Latin



In this chapter, you will learn

- Using advanced Pig Latin commands like JOIN, COGROUP, SAMPLE, SPLIT and STREAM
- Using operators and functions in Pig



Joining 2 or more data sets

- Pig Latin supports inner and outer joins of two or more relations
- Syntax for inner join:
`alias = JOIN alias BY field
[, alias BY field..]`
- Examples:
`joined = JOIN a BY f1, B BY f2;
joined = JOIN a BY $0, B BY $2, C BY $1;`



Pig Latin supports joining 2 or more data sets that have a field in common. Both inner and outer joins are supported. To perform an **inner** join, use the JOIN operator:
`alias = JOIN alias BY expression [, alias BY expression..]`

Example:

```
grunt> cat pets.txt
Doug      cat
Tom       dog
Mike      cat
Sarah     fish
grunt> cat hobbies.txt
Doug      reading
Tom       swimming
Mike      biking
Sarah     singing
grunt> pets = LOAD 'pets.txt';
grunt> hobbies = LOAD 'hobbies.txt';
grunt> joined = JOIN pets BY $0, hobbies BY $0;
grunt> DUMP joined;
(Tom,dog,Tom,swimming)
(Doug,cat,Doug,reading)
(Mike,cat,Mike,biking)
(Sarah,fish,Sarah,singing)
```

Note: Only equi-joins are supported.

Outer joins

- Pig can perform left, right or full outer joins (similar to SQL)
- The relation that has non-matching data must have a defined schema

- Syntax:

```
alias = JOIN alias BY field  
        [LEFT|RIGHT|FULL], alias BY field;
```



In the inner join example, everyone had a hobby and a pet. However, let's say Philip has only a hobby:

Philip reading

Philip would not be returned in an inner join. However, we can use an outer join to include the non-matching records from one or both relations.

Example:

```
grunt> hobbies = LOAD 'hobbies.txt';  
grunt> pets = load 'pets.txt' AS (name:chararray,  
hobby:chararray);  
grunt> joined = JOIN pets BY name RIGHT, hobbies BY $0;  
grunt> DUMP joined;
```

```
(Tom,dog,Tom,swimming)  
(Doug,cat,Doug,reading)  
(Mike,cat,Mike,biking)  
(Sarah,fish,Sarah,singing)  
(,,Philip,reading)
```

Note: it is necessary to provide a schema for pets because it needs to produce nulls for the non-matching records.

GROUP + JOIN = COGROUP

- COGROUP is similar to GROUP except multiple relations can be involved
- Relations are implicitly grouped on join field
- Syntax:

```
alias = COGROUP alias BY field,  
                alias BY field;
```
- Output is a set of tuples for each group key:
(group, {bag of records}, {bag of records})

records from first relation

records from second relation



COGROUP is a generalization of GROUP that can involve more than 1 relation.

The syntax is:

```
alias3 = COGROUP alias1 BY field [INNER], alias2 BY field  
[INNER];
```

The relations (alias1 and alias2) will be joined and grouped on the field they have in common. By default this is a **full outer join**, but the keyword `INNER` can be included for one or both relations.

The result is a relation where each record is (group key, bag-of-records, bag-of-records). The bags contain records from one of the input relations that match this group.

Although you can use `COGROUP` in place of a regular `GROUP`, it is a good idea to use `GROUP` when only one relation is used and `COGROUP` for multiple relations.

Example of COGROUP

- **pets.txt:**
Doug cat
Tom dog
Mike cat
- **hobbies.txt:**
Doug reading
Tom swimming
Mike biking
Philip reading
- `grpds = COGROUP pets BY $0, hobbies BY $0;`
`(Tom, {(Tom,dog)}, {(Tom,swimming)})`
`(Doug, {(Doug,cat)}, {(Doug,reading)})`
`(Mike, {(Mike,cat)}, {(Mike,biking)})`
`(Philip, {}, {(Philip,reading)})`



Take these two files as input:

pets.txt

Doug cat
Tom dog
Mike cat

hobbies.txt:


Doug reading
Tom swimming
Mike biking
Philip reading

If these files are COGROUPED on the person's name (field \$0) the output would be (formatting changed for readability):

(Tom,	{(Tom,dog)},	{(Tom,swimming)})
(Doug,	{(Doug,cat)},	{(Doug,reading)})
(Mike,	{(Mike,cat)},	{(Mike,biking)})
(Philip,	{},	{(Philip,reading)})

SAMPLE

- Use `SAMPLE` to choose a random set of tuples from a data set
- Syntax:
`alias = SAMPLE alias N;`



N should be a number
between 0-1,
for example .05



Sometimes it is useful to select a small subset of rows from a data set. `SAMPLE` does just that. When specifying the size of the relation, use a number between zero and one (such as .05 for 5% of the data).

Note that the input relation needs to be fully read in order to create the `SAMPLE`.

SPLIT

- A relation can be partitioned into 2 or more relations using SPLIT
- Syntax:
SPLIT alias INTO alias IF expression,
alias IF expression [, ...]

- Examples:

```
SPLIT users INTO males IF gender=='M', females  
IF gender=='F';
```

```
SPLIT a INTO b IF f1=='foo', c IF (f2<5 AND  
f3=='bar');
```



SPLIT is a useful command for dividing a relation into 2 or more relations. The input relation only needs to be scanned once to create the output data sets. The syntax is:

```
SPLIT alias INTO alias IF expression, alias IF expression  
[, alias IF expression..];
```

The expression can be simple comparison operator (e.g., `f1=='foo'`) or a compound expression that uses AND/OR. For compound expressions, enclose the expression in parenthesis.

The resulting relations can contain the same or different records from the input relation. For example:

```
SPLIT users INTO males IF gender=='M', engineers IF  
occupation=='engineer';
```

In the above statement, users becomes two relations: males and engineers. Some males are also engineers and would be outputted to both relations.

STREAM

- The `STREAM` operator sends a relation through an external script
- Examples:

```
b = STREAM a THROUGH `script.py`;
```

```
b = STREAM a THROUGH `cut -f 2`;
```



Like Hadoop streaming, Pig can send a data set to a script or program. The script reads the incoming records as tab-delimited lines and should return lines of tab-delimited fields. Common UNIX utilities can also be used, such as `cut`.

This example extracts the second field of the records in `a`:

```
b = STREAM a THROUGH `cut -f 2`;
```

Optionally, you can `DEFINE` a command for the script (especially useful if the command will be reused):

```
DEFINE mycmd `script.py`;  
b = STREAM a THROUGH mycmd;
```

Operators in Pig Latin

- Arithmetic:
 - `+` `-` `*` `/` `%` `?:`
- Comparison:
 - `==` `!=` `<` `>` `<=` `>=` `matches`
- NULL:
 - `IS NULL`, `IS NOT NULL`
- Boolean
 - `AND`, `OR`, `NOT`
- Others:
 - `FLATTEN`, cast operator



Pig Latin supports several operators very similar to most programming languages.

Arithmetic operators:

`+` (addition)
`-` (subtraction)
`*` (multiplication)
`/` (division)
`%` (modulo)
`?` (condition ? if-true : if-false), for example, `(name=='Doug' ? 'Found Doug' : 'Not Doug')`

Comparison operators:

`==` (equal)
`!=` (not equal)
`<` (less than), `>` (greater than)
`<=` (less than or equal), `>=` (greater than or equal)
`matches` (regular expression matching, using Java regex format)

Others:

`IS [NOT] NULL` (for NULL comparisons)
`AND`, `OR`, `NOT` (compound statements)
`FLATTEN` (see next page)
cast operator (change or identify a data type), for example: `(int)$1`

FLATTEN

- Unnest a bag or tuple

- Example:

```
(a, (b, c)) => GENERATE $0, FLATTEN($1) => (a, b, c)
```

- Removing extra bags:

- `((The), (cat), (in), (the), (hat.))`

- After FLATTEN..

- `(The)`
`(cat)`
`(in)`
`(the)`
`(hat.)`



FLATTEN is used to remove a level of nesting from a bag or tuples. For example:
(a, (b,c)) can be "flattened" into (a,b,c)

FLATTEN is often used to unnest the result of a function. For example:

```
> cat data.txt;
The cat in the hat.
> lines = LOAD 'data.txt' USING TextLoader();
> DUMP lines:
(The cat in the hat.)
> words = FOREACH lines GENERATE TOKENIZE($0);
> DUMP words;
((The), (cat), (in), (the), (hat.))
> flat = FOREACH words GENERATE FLATTEN($0);
> DUMP flat;
(The)
(cat)
(in)
(the)
(hat.)
```

Built-in Functions

- A few built-in functions:
 - AVG - average of the values in a column
 - CONCAT - concatenate 2 strings
 - COUNT - count the number of elements in a bag, ignore NULL
 - COUNT_STAR - count, including NULLs
 - DIFF - find the differing elements
 - IsEMPTY - Tests if a bag is empty
 - MAX/MIN - maximum/minimum value in a column
 - SIZE - the number of elements in a data set
 - SUM - add the values in a column
 - TOKENIZE - split a string into words



Pig comes with a few built-in functions. Many are aggregate functions such as AVG, COUNT, MAX, MIN and SUM. There are also functions for concatenating strings, find differences between elements, and tokenizing a string.

The aggregate functions require a previous GROUP statement. Use GROUP ALL for global calculations:

```
users = LOAD 'data';  
grp = GROUP users ALL;  
total = FOREACH grp GENERATE COUNT(users);
```

Remember that function names are case-sensitive.

Using UDFs

- Pig allows user-defined functions written in Java

```
public class MyFunc extends EvalFunc {  
    public Double exec(Tuple input) {  
        ...  
    }  
}  
  
grunt> REGISTER my-code.jar  
grunt> DEFINE myFunc com.examples.MyFunc();  
grunt> b = FOREACH a GENERATE myFunc($0);
```



One way of extending the Pig Latin language is by writing user-defined functions (UDFs). These currently need to be written in Java. The steps required are:

1. Write a class that extends `EvalFunc` and implement the `exec` method.
2. Compile and package into a jar
3. Tell pig about the jar using the `REGISTER` keyword
4. Optionally `DEFINE` a function name. Without this step, the fully-qualified class name is the function name (e.g., `com.examples.MyFunc()`)
5. Invoke your function in the pig script

PiggyBank

- A library of common functions written by the community called "PiggyBank"
- <http://wiki.apache.org/pig/PiggyBank>
- Common functions for math, parsing dates and strings and custom loaders



There is also a set of functions written by the community. It's called the PiggyBank. Details can be found here: <http://wiki.apache.org/pig/PiggyBank>

Dates:

- CustomFormatToISO - convert arbitrary date format to ISO format
- UnixToISO and ISOToUnix - convert between ISO format and Unix timestamps

Math:

- ABS - absolute value of a number
- LOG - natural log of a number
- POW - a number raised to a power
- RANDOM - return a random number
- ROUND - round numbers to the closest long

Strings:

- INDEXOF - search for a string
- LENGTH - the length of a string
- LOWER - convert to lowercase
- SUBSTRING - extract a portion of a string
- UPPER - convert to uppercase

Storage:

- MyRegexLoader - parse a file given a user-defined regular expression
- SequenceFileLoader - read Hadoop SequenceFile format
- XMLLoader - parses XML files by a user-supplied start/end tag

In this chapter, you have learned

- Using advanced Pig Latin commands like JOIN, COGROUP, SAMPLE, SPLIT and STREAM
- Using operators and functions in Pig

